

HEALTHCARE PROJECT

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import OrdinalEncoder
from sklearn.ensemble import RandomForestRegressor
import scipy.stats as stats
import sklearn.pipeline as Pipeline
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error as mse
from sklearn.model_selection import KFold, train_test_split
import plotly.express as px
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import SGDRegressor, Ridge
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
```

```
hospital=pd.read_csv(r"C:\Users\hp\AppData\Roaming\Microsoft\Windows\Start Menu\Hospitalisation details.csv")
medical=pd.read_csv(r"C:\Users\hp\AppData\Roaming\Microsoft\Windows\Start Menu\Medical Examinations.csv")
Names=pd.read_csv(r"C:\Users\hp\AppData\Roaming\Microsoft\Windows\Start Menu\Names1.csv")
```

hospital.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2343 entries, 0 to 2342
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Customer ID     2343 non-null  object
1   year            2343 non-null  object
2   month           2343 non-null  object
3   date            2343 non-null  int64
4   children        2343 non-null  int64
5   charges         2343 non-null  float64
6   Hospital tier   2343 non-null  object
7   City tier        2343 non-null  object
8   State ID        2343 non-null  object
dtypes: float64(1), int64(2), object(6)
memory usage: 164.9+ KB
```

medical.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Customer ID     2335 non-null  object
1   BMI             2335 non-null  float64
2   HBA1C           2335 non-null  float64
3   Heart Issues    2335 non-null  object
```

```
Names.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Customer ID 2335 non-null   object
 1   name        2335 non-null   object
dtypes: object(2)
memory usage: 36.6+ KB
```

```
hospital.isnull().sum()
```

```
Customer ID    0
year           0
month          0
date           0
children       0
charges        0
Hospital tier   0
City tier       0
State ID       0
dtype: int64
```

```
medical.isnull().sum()
```

```
Customer ID    0
BMI            0
HBA1C          0
Heart Issues   0
Any Transplants 0
Cancer history 0
NumberOfMajorSurgeries 0
smoker         0
dtype: int64
```

```
Names.isnull().sum()
```

```
Customer ID    0
name           0
dtype: int64
```

```
#there are no missing values in the data
#we can conclude hospital has some more entries than medical and names table
```

```
#now we will merge the data(note we can merge 2 table at once)
main_data=pd.merge(hospital,medical,how='inner',on='Customer ID')
```

```
main_data=main_data.merge(Names, how='inner',on='Customer ID')
```

```
#now we got all 3 tables combined in 1 table as main_data
main_data.describe()
```

	date	children	charges	BMI	HBA1C
count	2335.000000	2335.000000	2335.000000	2335.000000	2335.000000
mean	15.563597	1.025696	13529.918034	30.972649	6.578998
std	8.720508	1.234754	11898.654299	8.742095	2.228731
min	1.000000	0.000000	563.840000	15.010000	4.000000
25%	8.000000	0.000000	5084.010000	24.600000	4.900000
50%	15.000000	0.000000	9630.910000	30.400000	5.810000
75%	23.000000	2.000000	16912.295000	36.300000	7.955000
max	30.000000	5.000000	63770.430000	55.050000	12.000000

```
#note when mean and 50% of the data matches means the data is good and normalised data
```

```
main_data.head()
```

	Customer ID	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	NumberOfMajorSurgeries	smoker	name
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3	R1013	17.58	4.51	No	No	No	1	No	German, Mr. Aaron K
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1	R1013	17.60	4.39	No	No	No	1	No	Rosendahl, Mr. Evan P
2	Id2333	1993	Jun	30	0	600.00	tier - 2	tier - 1	R1013	16.47	6.35	No	No	Yes	1	No	Albano, Ms. Julie
3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3	R1013	17.70	6.28	No	No	No	1	No	Riveros Gonzalez, Mr. Juan D. Sr.
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3	R1013	22.34	5.57	No	No	No	1	No	Brietzke, Mr. Jordan

```
#to know the trivial values from the data we can find rows with '?'
```

```
(main_data=='?').sum()
```

```
Customer ID      0
year             2
month            3
date             0
children         0
charges          0
Hospital tier     1
City tier         1
State ID         2
BMI              0
HBA1C            0
Heart Issues     0
Any Transplants  0
Cancer history   0
NumberOfMajorSurgeries 0
smoker           2
name             0
dtype: int64
```

```
#we will calculate the % of data having trivial value
```

```
miss_value=(main_data=='?').sum(axis=1)/main_data.shape[1]*100
```

```
miss_value[miss_value>0]
```

```
miss_value[miss_value>0]
```

```
11      5.882353
13      5.882353
17     11.764706
542     5.882353
1046    5.882353
1049    5.882353
1700    5.882353
1775    5.882353
2165    5.882353
2332    5.882353
dtype: float64
```

```
#as we can see above there are 10 rows
main_data.shape
```

```
(2335, 17)
```

```
#we are deleting the rows with missing value
```

```
Data=main_data.drop(index=miss_value[miss_value>0].index)
```

```
#we are deleting the rows with missing value
Data=main_data.drop(index=miss_value[miss_value>0].index)
```

```
Data.shape
```

```
(2325, 17)
```

```
#rows got deleted successfully
```

```
Data.columns
```

```
Index(['Customer ID', 'year', 'month', 'date', 'children', 'charges',  
       'Hospital tier', 'City tier', 'State ID', 'BMI', 'HBA1C',  
       'Heart Issues', 'Any Transplants', 'Cancer history',  
       'NumberOfMajorSurgeries', 'smoker', 'name'],  
      dtype='object')
```

```
#Nominal value- values that are not numeric like color,  
                #(heart issues,any transplant, cancer history, smoker,state id)
```

```
#Ordinal value- values that are clearly shows order and rank  
                #(Hospital tier, city tier)
```

```
Group=Data[["Hospital tier","City tier"]]
```

```
Group
```

```
[25]:
```

	Hospital tier	City tier
--	---------------	-----------

0	tier - 2	tier - 3
1	tier - 2	tier - 1
2	tier - 2	tier - 1
3	tier - 3	tier - 3
4	tier - 3	tier - 3
...
2329	tier - 1	tier - 3
2330	tier - 1	tier - 2
2331	tier - 1	tier - 3
2333	tier - 2	tier - 3
2334	tier - 1	tier - 3

2325 rows × 2 columns

```
[26]: ordinal = OrdinalEncoder(categories= [['tier - 3', 'tier - 2', 'tier - 1'], ['tier - 3', 'tier - 2', 'tier - 1']])
Data[['city_order', 'hospital_order']] = ordinal.fit_transform(Data[['City tier', 'Hospital tier']])
```

```
[27]: Data
```

```
[27]:
```

	Customer ID	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	NumberOfMajorSurgeries	smoker	name
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3	R1013	17.580	4.51	No	No	No	1	No	Germ Mr. Aar
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1	R1013	17.600	4.39	No	No	No	1	No	Rosendi Mr. Eval
2	Id2333	1993	Jun	30	0	600.00	tier - 2	tier - 1	R1013	16.470	6.35	No	No	Yes	1	No	Alba Ms. Ju
3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3	R1013	17.700	6.28	No	No	No	1	No	Rive Gonzal Mr. Ju D.
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3	R1013	22.340	5.57	No	No	No	1	No	Brietz Mr. Jorc
...
2329	Id6	1962	Aug	4	0	52590.83	tier - 1	tier - 2	R1011	32.800	6.59	No	No	No	No major surgery	yes	Baker, I Rurrel

```
[28]: #now 2 new dummy columns got created
```

```
[29]: # now we will check the state ID
Data[["State ID"]].value_counts()
```

```
[29]: State ID
R1013      609
R1011      574
R1012      572
R1024      159
R1026       84
R1021       70
R1016       64
R1025       40
R1023       38
R1017       36
R1019       26
R1022       14
R1014       13
R1015       11
R1018        9
R1020        6
Name: count, dtype: int64
```

```
[30]: #now we have 3 major states but all other has low count if we make dummies for all that it will hamper the decision
      #so we will combine all other stateId to one new state Id
```

```
[31]: varid=Data[["State ID"]].value_counts()
```

```
[32]: varid[0:3]
```

```
[32]: State ID
      R1013      609
      R1011      574
      R1012      572
      Name: count, dtype: int64
```

```
[33]: Data.columns = Data.columns.str.replace(' ', '_')
      Data.columns
```

```
[33]: Index(['Customer_ID', 'year', 'month', 'date', 'children', 'charges',
          'Hospital_tier', 'City_tier', 'State_ID', 'BMI', 'HBA1C',
          'Heart_Issues', 'Any_Transplants', 'Cancer_history',
          'NumberOfMajorSurgeries', 'smoker', 'name', 'city_order',
          'hospital_order'],
          dtype='object')
```

```
#Now we need to find,str values instead of numbers
```

```
Data.NumberOfMajorSurgeries.unique()
```

```
array(['1', 'No major surgery', '2', '3'], dtype=object)
```

```
Data.loc[Data.NumberOfMajorSurgeries == 'No major surgery'] = 0
```

```
Data.NumberOfMajorSurgeries.unique()
```

```
array(['1', 0, '2', '3'], dtype=object)
```

```
Data.year.unique()
```

```
array(['1992', '1993', '1998', 0, '1995', '1997', '2004', '2000', '2003',
      '1988', '1987', '1986', '1984', '1983', '1979', '1973', '1972',
      '1975', '1970', '1969', '1966', '1963', '1964', '1961', '1959',
      '1958'], dtype=object)
```

```
#we need to calculate the age based on year and make a new column of that
Data.year = Data.year.astype(int)
Data['age']=2025 - Data.year
Data.age.unique()
```

```
array([ 33,  32,  27, 2025,  30,  28,  21,  25,  22,  37,  38,
        39,  41,  42,  46,  52,  53,  50,  55,  56,  59,  62,
        61,  64,  66,  67])
```

```
[40]: #as we have 0 in sum of the year blocks so we need to make it as 0
Data.loc[Data.age == 2025] = 0

[41]: Data.age.unique()

[41]: array([33, 32, 27,  0, 30, 28, 21, 25, 22, 37, 38, 39, 41, 42, 46, 52, 5
        50, 55, 56, 59, 62, 61, 64, 66, 67])

[42]: #we need to make gender column
Data.name.head()

[42]: 0          German, Mr.  Aaron K
1      Rosendahl, Mr.  Evan P
2          Albano, Ms.  Julie
3  Riveros Gonzalez, Mr.  Juan D. Sr.
4      Brietzke, Mr.  Jordan
Name: name, dtype: object

[43]: Data['salutation'] = Data.name.str.split('[,.]').str[1] #1

[44]: Data.salutation.unique()

[44]: array([' Mr', ' Ms', nan, ' Mrs'], dtype=object)

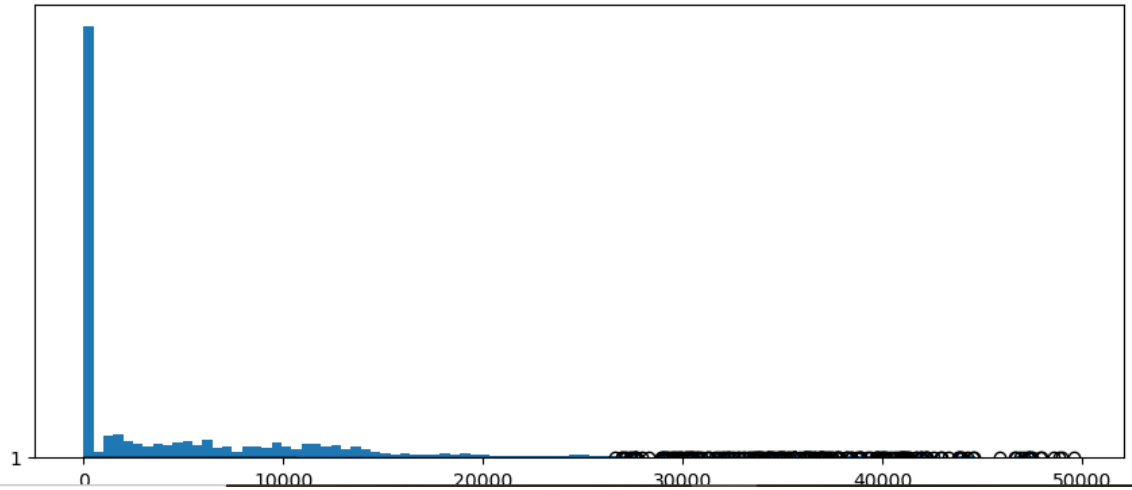
[45]: #Data.rename(columns={'salutation': 'Gender'}, inplace=True)
```

```
[46]: Data['gender'] = 'female'
Data.loc[Data.salutation == 'Mr', 'gender'] = 'male'


[47]: Data.loc[Data.salutation == 'Mrs']
```



```
[49]: plt.figure(figsize=(25,10))
grid = plt.GridSpec(2, 2, wspace=0.4, hspace=0.3)
plt.subplot(grid[0, 0])
plt.hist(Data.charges, bins = 100)
plt.boxplot(Data.charges, vert = False)
plt.show()
```




```
plt.figure(figsize = (25,10))
sns.boxplot(x='charges',y='City_tier', data=Data)
plt.show()
```



```
[51]: plt.figure(figsize=(25,10))
sns.boxplot(x='charges',y='Hospital_tier', data=Data)
plt.show
```

```
[51]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[52]: #now we need to find the median Hospital considering all tier of Hospitals using radar chart
Median=Data.groupby('Hospital_tier')[['charges']].median().reset_index()
```

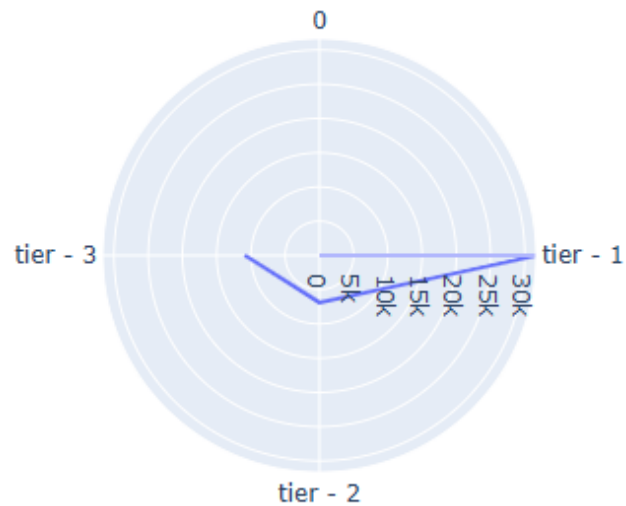
```
[53]: Median
```

```
[53]:
```

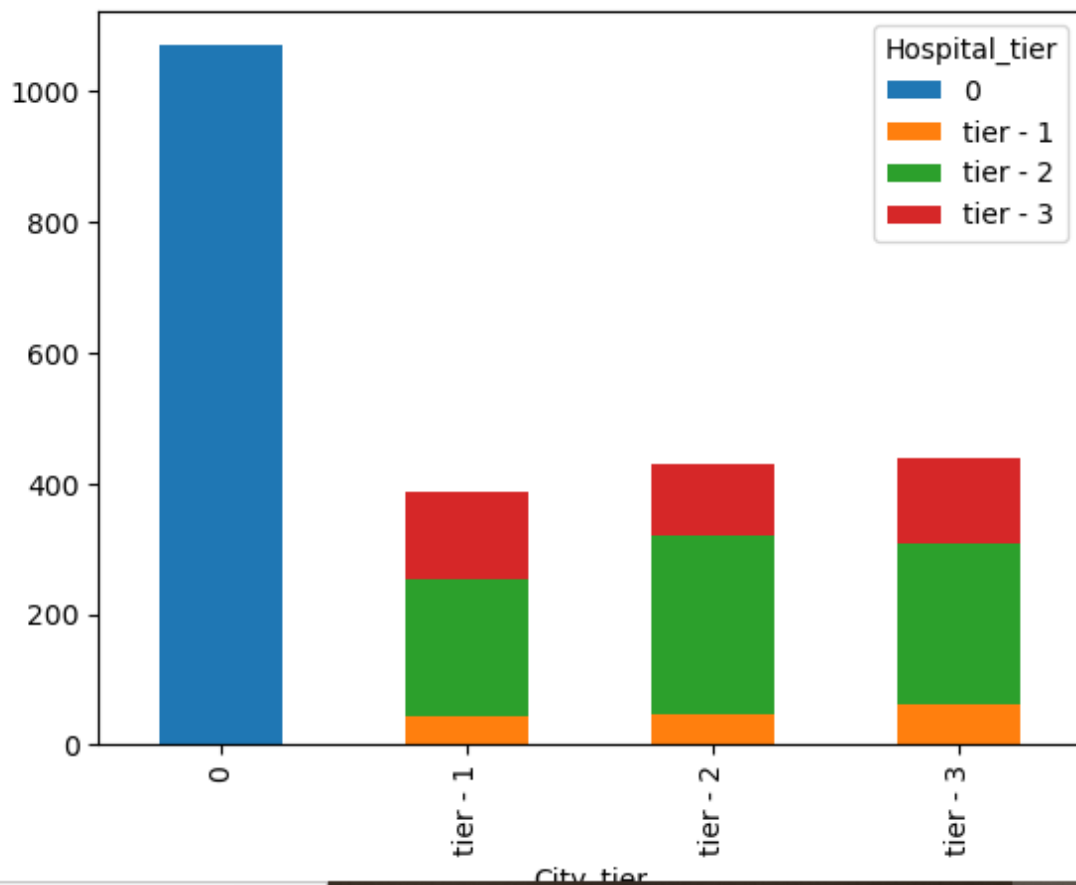
	Hospital_tier	charges
0	0	0.000
1	tier - 1	31660.385
2	tier - 2	6868.390
3	tier - 3	10874.485

```
[54]: Radar=px.line_polar(Median,r='charges',theta='Hospital_tier')
```

```
[55]: Radar.show()
```



```
pd.crosstab(Data.City_tier ,Data.Hospital_tier).plot.bar(stacked=True)  
plt.show()
```



```
57]: pd.crosstab(Data.City_tier, Data.Hospital_tier)
```

```
57]: Hospital_tier    0  tier - 1  tier - 2  tier - 3
```

City_tier					
	0				
	1070	0	0	0	
tier - 1	0	44	209	133	
tier - 2	0	47	273	111	
tier - 3	0	63	245	130	

```
58]: #Now we will check for the null hypothesis  
#Average hospitalization cost across the 3 types of hospitals is not sig
```

```
59]: Data.columns
```

```
59]: Index(['Customer_ID', 'year', 'month', 'date', 'children', 'charges',  
          'Hospital_tier', 'City_tier', 'State_ID', 'BMI', 'HBA1C',  
          'Heart_Issues', 'Any_Transplants', 'Cancer_history',  
          'NumberOfMajorSurgeries', 'smoker', 'name', 'city_order',  
          'hospital_order', 'age', 'salutation', 'gender'],  
          dtype='object')
```

```
[60]: Data[['charges', 'Hospital_tier']].tail(50)
```

```
[60]:
```

	charges	Hospital_tier
2284	43753.34	tier - 2
2285	43813.87	tier - 2
2286	43817.45	tier - 2
2287	0.00	0
2288	43921.18	tier - 2
2289	43943.88	tier - 2
2290	44202.65	tier - 2
2291	0.00	0
2292	0.00	0
2293	0.00	0
2294	44501.40	tier - 2
2295	44585.46	tier - 2
2296	0.00	0
2297	0.00	0

```
[61]: hospital=Data.groupby(['Hospital_tier']).charges.mean()
```

```
[62]: print(hospital)
```

```
Hospital_tier
0          0.000000
tier - 1    29708.588052
tier - 2    11874.593040
tier - 3     9497.587701
Name: charges, dtype: float64
```

```
[63]: from statsmodels.formula.api import ols
import statsmodels.api as sm
```

```
[64]: mod = ols('charges ~ Hospital_tier', data = Data).fit()
      res = sm.stats.anova_lm(mod)
      res
```

```
[64]:
```

	df	sum_sq	mean_sq	F	PR(>F)
Hospital_tier	3.0	1.513512e+11	5.045041e+10	938.988595	0.0
Residual	2321.0	1.247038e+11	5.372846e+07	NaN	NaN

```
[65]: #so as we can see above P-value the average charges taken by 3 different categories of hospital
      #has significant difference therefore we reject the null hypothesis
```

```
[66]: #Average hospitalization cost across the 3 types of cities is not significantly different
city=Data.groupby(['City_tier']).charges.mean()
```

```
[67]: print(city)

City_tier
0          0.000000
tier - 1    12998.467668
tier - 2    13096.768910
tier - 3    13922.222785
Name: charges, dtype: float64
```

```
[68]: mod = ols('charges ~ City_tier', data = Data).fit()
      res = sm.stats.anova_lm(mod)
      res
```

```
[68]:
```

	df	sum_sq	mean_sq	F	PR(>F)
City_tier	3.0	1.032260e+11	3.440865e+10	462.089465	2.240926e-235
Residual	2321.0	1.728290e+11	7.446318e+07	NaN	NaN

```
[69]: #as we can see average hospitalization cost is not significantly different
      #so we fail to reject the null hypothesis
```

```
[70]: # Average hospitalization cost for smokers is not significantly different than non-smokers
```

```
[71]: smoker = Data.loc[Data.smoker == 'yes', 'charges']
      no_smoker=Data.loc[Data.smoker != 'yes', 'charges']
```

```
[72]: stats.ttest_ind(smoker, no_smoker)
```

```
[72]: TtestResult(statistic=69.96251699445597, pvalue=0.0, df=2323.0)
```

```
[73]: #from p-value smoker charges to non-smoker charges are different so we reject the null hypothesis
```

```
[74]: #Smoking and heart issues are independant
```

```
[75]: table = pd.crosstab(Data.smoker, Data.Heart_Issues)
      table
```

```
[75]: Heart_Issues    0  No  yes
```

smoker				
0	1070	0	0	
No	0	493	493	
yes	0	136	133	

```
[76]: chi, p, df, expected = stats.chi2_contingency(table)
```

```
[77]: chi, p, df, expected
```

```
[77]: (2325.0486973279203,
       0.0,
       4,
       array([[492.43010753, 289.47526882, 288.09462366],
              [453.77204301, 266.75010753, 265.47784946],
              [123.79784946, 72.77462366, 72.42752688]]))
```

```
[78]: #Looking at the p_value, we fail to reject the null hypothesis
```

```
[79]: #Check the correlation between predictors to identify highly correlated predictors. Visualize using a heatmap.  
#important predicting columns are-  
Data.columns
```

```
[79]: Index(['Customer_ID', 'year', 'month', 'date', 'children', 'charges',
        'Hospital_tier', 'City_tier', 'State_ID', 'BMI', 'HBA1C',
        'Heart_Issues', 'Any_Transplants', 'Cancer_history',
        'NumberOfMajorSurgeries', 'smoker', 'name', 'city_order',
        'hospital_order', 'age', 'salutation', 'gender'],
        dtype='object')
```


[illegible]

```
[81]: new_data
```

[illegible]

```
[82]: numeric_data = new_data.select_dtypes(exclude='object').corr()
      ma = np.ones_like(numeric_data)
```

```
[83]: plt.figure(figsize = (18,5))
      sns.heatmap(numeric_data, annot= True , cmap='PuRd')
      plt.show()
```



```
[84]: #5fold cross validation
      data_2 = pd.get_dummies(new_data, drop_first=True)
      data_2.reset_index(drop=True, inplace = True)
```

```
[85]: data_2.head()
```

```
[85]:
```

	children	charges	BMI	HBA1C	age	Heart_Issues_No	Heart_Issues_yes	Any_Transplants_No	Any_Transplants_yes	Cancer_history_No	Cancer_history_Yes	NumberOfMajorSurgeries
0	0	563.84	17.58	4.51	33	True	False	True	False	True	False	0
1	0	570.62	17.60	4.39	33	True	False	True	False	True	False	0
2	0	600.00	16.47	6.35	32	True	False	True	False	False	True	0
3	0	604.54	17.70	6.28	33	True	False	True	False	True	False	0
4	0	637.26	22.34	5.57	27	True	False	True	False	True	False	0

```
[86]: model_data = data_2.drop(columns = 'charges')
      model_data.head()
      model_data['charges'] = data_2.charges
      model_data.head()
```

```
[86]:
```

	children	BMI	HBA1C	age	Heart_Issues_No	Heart_Issues_yes	Any_Transplants_No	Any_Transplants_yes	Cancer_history_No	Cancer_history_Yes	NumberOfMajorSurgeries
0	0	17.58	4.51	33	True	False	True	False	True	False	0
1	0	17.60	4.39	33	True	False	True	False	True	False	0
2	0	16.47	6.35	32	True	False	True	False	False	True	0
3	0	17.70	6.28	33	True	False	True	False	True	False	0
4	0	22.34	5.57	27	True	False	True	False	True	False	0

```
[87]: X=data_2.drop(columns = 'charges')
      y=data_2['charges']
```

```
[88]: pipeline = Pipeline(steps=[('scaler', StandardScaler()), ('regressor', Ridge())])
```

```
[89]: parameters = {'regressor__alpha': [0.001, 0.01, 0.1, 1, 10, 100]}
```

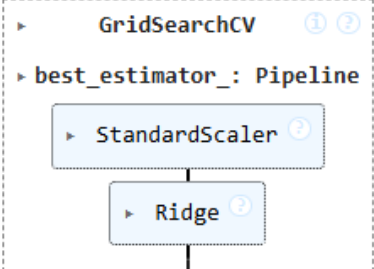


```
[90]: kfold = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[91]: model_ridge = GridSearchCV(pipeline, parameters, cv=kfold, scoring='neg_mean_squared_error')
```

```
[92]: model_ridge.fit(X, y)
```

```
[92]:
```



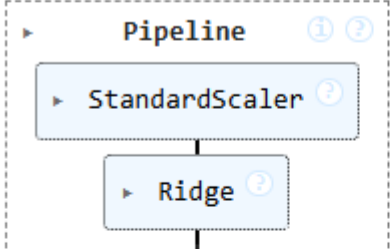
The diagram shows a `GridSearchCV` object (labeled 1 and 2) containing a `best_estimator_` attribute (labeled 3) which is a `Pipeline` object. The `Pipeline` contains a `StandardScaler` (labeled 2) and a `Ridge` (labeled 2) estimator.

```
[93]: model_ridge.best_params_
```

```
[93]: {'regressor__alpha': 1}
```

```
[94]: model_ridge.best_estimator_
```

```
[94]:
```



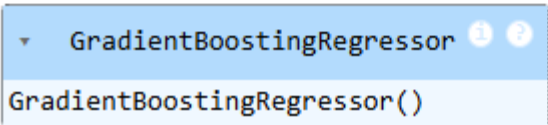
The diagram shows a `Pipeline` object (labeled 1 and 2) containing a `StandardScaler` (labeled 2) and a `Ridge` (labeled 2) estimator.

```
[97]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[98]: X_train,X_test,y_train,y_test = train_test_split(X,y)
```

```
[99]: model = GradientBoostingRegressor()  
model.fit(X_train, y_train)
```

```
[99]:
```



The diagram shows a `GradientBoostingRegressor` object (labeled 1 and 2) with the `GradientBoostingRegressor()` constructor.

```
[100]: print(model.feature_importances_)
```

```
[9.26529561e-03 1.18484675e-01 1.17478367e-02 1.69330114e-01
 1.86651393e-04 0.00000000e+00 7.93831112e-05 1.98567593e-04
 6.47220664e-07 1.64507916e-04 2.38864174e-05 4.69668547e-05
 7.23706261e-06 5.85556072e-04 6.89878675e-01]
```

```
[101]: model.score(X_train,y_train)
```

```
[101]: 0.9528373650472055
```

```
[102]: model.score(X_test,y_test)
```

```
[102]: 0.9516113339047438
```

```
[103]: model_data.columns
```

```
[103]: Index(['children', 'BMI', 'HBA1C', 'age', 'Heart_Issues_No',
        'Heart_Issues_yes', 'Any_Transplants_No', 'Any_Transplants_yes',
        'Cancer_history_No', 'Cancer_history_Yes', 'NumberOfMajorSurgeries_1',
        'NumberOfMajorSurgeries_2', 'NumberOfMajorSurgeries_3', 'smoker_No',
        'smoker_yes', 'charges'],
        dtype='object')
```

```
[118]: pred_data = pd.DataFrame({'Name' : ['Christopher, Ms. Jayna'],
        'DOB' : ['12/28/1988'],
        'city_tier' : ['tier - 1'], 'children' : [2],
        'HbA1c' : [5.8],
        'smoker_yes' : [1],
        'heart_issues_yes' : [0],
        'any_transplants_yes' : [0],
        'numberofmajorsurgeries' : [0],
        'cancer_history_yes' : [1],
        'hospital_tier' : ['tier - 1'],
        'bmi' : [85/(1.70 **2)],
        'state_id_R1011' : [1]
        })
```

```
[119]: pred_data
```

```
[119]:
```

	Name	DOB	city_tier	children	HbA1c	smoker_yes	heart_issues_yes	any_transplants_yes	numberofmajorsurgeries	cancer_history_yes	hospital_tier	
0	Christopher, Ms. Jayna	12/28/1988	tier - 1	2	5.8	1	0	0	0	1	tier - 1	29.41

```
[120]: pred_data['gender_male'] = 0
pred_data.loc[pred_data.Name.str.split(',').str[1] == 'Mr', 'gender_male'] = 1
pred_data.drop(columns = 'Name', inplace = True)
```

```
[120]: pred_data['gender_male'] = 0
pred_data.loc[pred_data.Name.str.split(',').str[1] == 'Mr', 'gender_male'] = 1
pred_data.drop(columns = 'Name', inplace = True)
```

```
[109]:
```

```
[109]:
```

	dob	city_tier	children	hba1c	smoker_yes	heart_issues_yes	any_transplants_yes	numberofmajorsurgeries	cancer_history_yes	hospital_tier	bmi	state
0	12/28/1988	tier - 1	2	5.8	1	0	0	0	1	tier - 1	29.411765	

```
[122]: pred_data.drop(columns = 'DOB', inplace = True)
```

```
[123]: pred_data[['city_tier_ord', 'hospital_tier_ord']] = ordinal.transform(pred_data[['city_tier', 'hospital_tier']])
```

```
[124]: pred_data.drop(columns=['city_tier', 'hospital_tier'], inplace = True )
```

```
[125]: for col in model_data.columns:
        if col not in pred_data.columns and col != 'charges':
            pred_data[col] = 0
```

```
[126]: pred_data
```

	children	HbA1c	smoker_yes	heart_issues_yes	any_transplants_yes	numberofmajorsurgeries	cancer_history_yes	bmi	state_id_R1011	gender_male	...	Heart_Is
0	2	5.8	1	0	0	0	1	29.411765	1	0	...	

1 rows × 23 columns

```
[127]: model_data.columns
```

```
[127]: Index(['children', 'BMI', 'HBA1C', 'age', 'Heart_Issues_No',
        'Heart_Issues_yes', 'Any_Transplants_No', 'Any_Transplants_yes',
        'Cancer_history_No', 'Cancer_history_Yes', 'NumberOfMajorSurgeries_1',
        'NumberOfMajorSurgeries_2', 'NumberOfMajorSurgeries_3', 'smoker_No',
        'smoker_yes', 'charges'],
        dtype='object')
```

```
[127]: model_data.columns
```

```
[127]: Index(['children', 'BMI', 'HBA1C', 'age', 'Heart_Issues_No',
        'Heart_Issues_yes', 'Any_Transplants_No', 'Any_Transplants_yes',
        'Cancer_history_No', 'Cancer_history_Yes', 'NumberOfMajorSurgeries_1',
        'NumberOfMajorSurgeries_2', 'NumberOfMajorSurgeries_3', 'smoker_No',
        'smoker_yes', 'charges'],
        dtype='object')
```

```
[128]: pred_data.columns
```

```
[128]: Index(['children', 'HbA1c', 'smoker_yes', 'heart_issues_yes',
        'any_transplants_yes', 'numberofmajorsurgeries', 'cancer_history_yes',
        'bmi', 'state_id_R1011', 'gender_male', 'BMI', 'HBA1C', 'age',
        'Heart_Issues_No', 'Heart_Issues_yes', 'Any_Transplants_No',
        'Any_Transplants_yes', 'Cancer_history_No', 'Cancer_history_Yes',
        'NumberOfMajorSurgeries_1', 'NumberOfMajorSurgeries_2',
        'NumberOfMajorSurgeries_3', 'smoker_No'],
        dtype='object')
```

```
[129]: pred_data=pred_data[model_data.drop(columns='charges').columns]
```

```
[130]: model.predict(pred_data)
```

```
[130]: array([22627.63241263])
```

END